


ECE4078 Intelligent robotics



Lab 3-2: Kal Backman & Shujie Zhou

ECE4078 Intelligent robotics

Week	Objectives	Milestones
2: M1-1	Introduction and setup	
3: M2-1	Calibration, ARUCO markers	M1 due
4: M2-2	SLAM	
5: M2-3	SLAM	
6: M3-1	Object recognition & localisation	M2 due
* 7: M3-2	Object recognition & localisation	
8: M4-1	Navigation & Planning	M3 due
9: M4-2	Navigation & Planning	
10: M5-1	Integration	M4 due
11: Final	Improvement	M5 due
12: Final	Final demo & competition	Final demo

Milestone 3

Milestone is focused on object recognition and localisation

Your task is to classify and estimate the position of various fruits given a single image

Task 1: Gather training data (Week 6)

Use the provided scripts to generate data for the default segmentation network or create your own data generator for YOLO

Task 2: Training the network (Week 6)


Train the default network using the “main.py” script in the “network/scripts/” folder or your own YOLO network from your chosen GitHub repository

Task 3: Estimating object poses (Week 7)

Given your trained network’s output, determine the physical position of the observed object

Estimating object pose


Given any standard RGB image, can we calculate the 3D position of any pixel within an image?



Estimating object pose


What if we had 2 RGB images?

We can triangulate the projected light rays to figure out the 3D position of the corresponding point




Estimating object pose

This is difficult to do however as it is reliant on us being confident on the relative poses of both cameras and choosing a pixel in both images that correspond to the exact same spot on the human



Estimating object pose

What if we project two lines from the camera?




If the human were to walk backwards, away from the camera, they should appear smaller in the image

Estimating object pose




This will ONLY work if:

We know the height of the object we are looking at

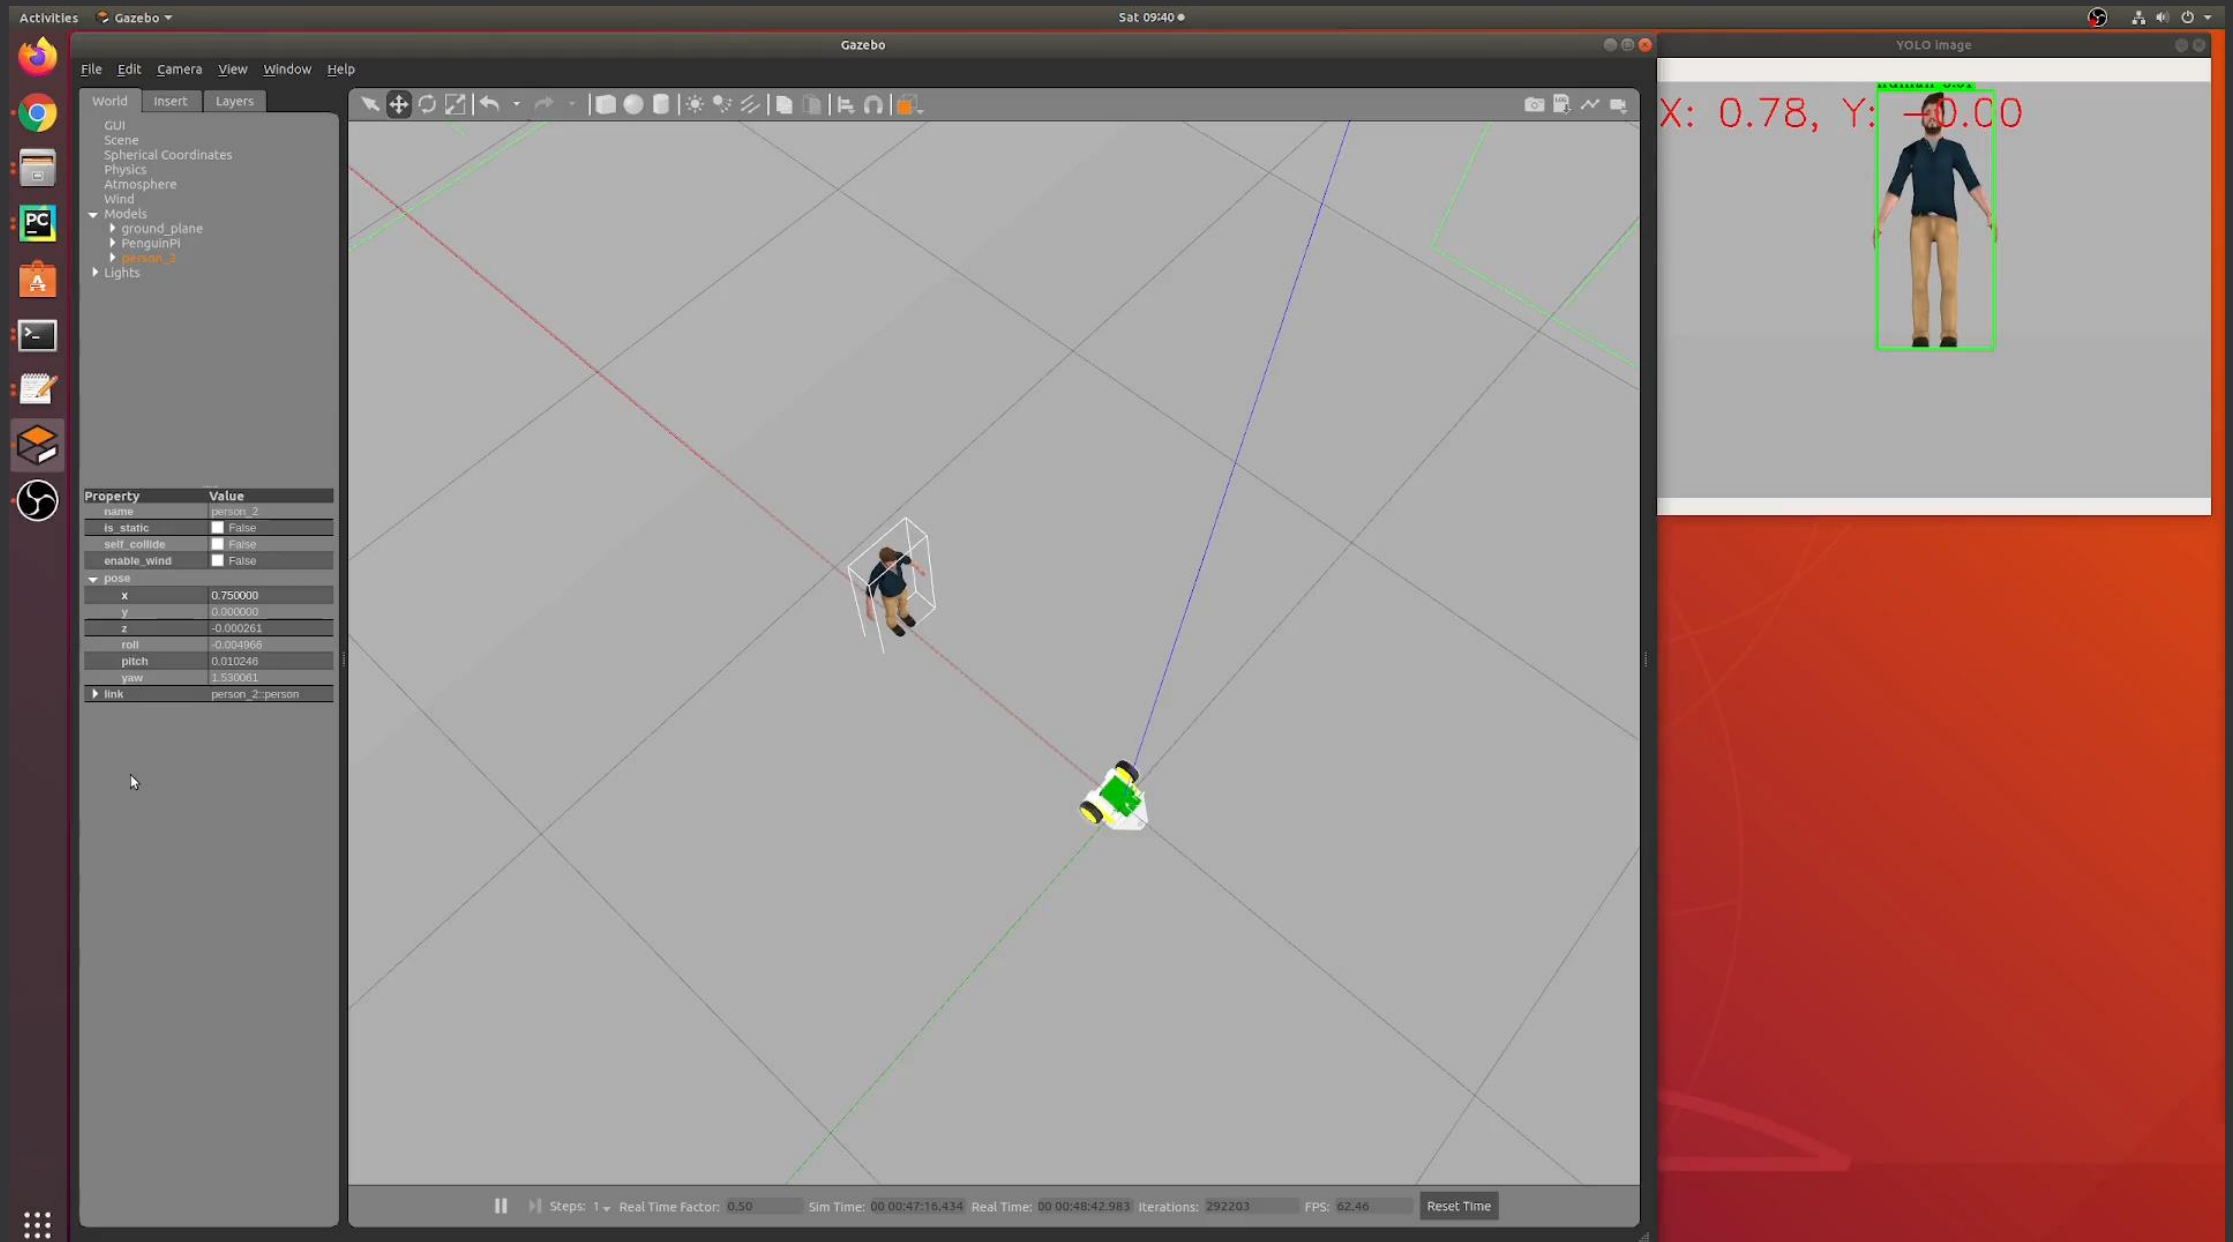
The object's height does not change



Estimating object pose



Estimating object pose



Estimating object pose

Things you will need to calculate fruit pose:

- Camera's intrinsic parameters, which you would have calculated in M2

$$\begin{bmatrix} F_x & 0 & u_0 \\ 0 & F_y & v_0 \\ 0 & 0 & 1 \end{bmatrix}$$

- Height in pixels and x-pixel coordinate of the object
 - YOLO provides this information to you with bounding boxes
 - The default network requires you to create bounding boxes from your segmentation map
Take a look at OpenCV's "findContours" or "SimpleBlobDetector" to do this

- Physical height of the objects

The physical heights of the simulated fruits can be found in "TargetPoseEst.py" or you can measure them in gazebo. For the physical fruit you can measure them with a ruler.

You can also use the physical width of the fruit as an additional estimate (make sure the fruit has a constant width from all viewing angles)

Make sure if your network is trained on a different image size compared to the PenguinPi camera resolution of 640x480, then you scale your pixel coordinates

Eg: Network trained on 320x240 images will need all pixel values scaled by 2x

Milestone 3

We will mark you based on the output of CV_eval.py script which compares the true global positions of the fruit with that of your global fruit estimates

You will have to modify TargetPoseEst.py to create your global fruit estimates

TargetPoseEst.py will receive a list of images along with their associated poses they were taken at as input

The input images will be taken as if someone was mapping out the arena like in M2

There are only 2 of each fruit in the marking map, the same fruit may appear in multiple images therefore you will have to perform some sort of filter / clustering algorithm to ONLY output 2 of each fruit

You will be marked on both a simulated dataset and a physical dataset

You can modify TargetPoseEst.py script as much as you like (even completely rewrite it, which I recommend) as long as your script adheres to the same input and output format as the skeleton code

In your submission make sure you submit:

- TargetPoseEst.py - To generate your global fruit map (you can upload 2 versions, sim & physical if you like)

- Network related scripts - scripts used for inference (NOT training scripts)

- Model weights - If weights are too large for moodle you can submit a link to a google drive file

- Read me file - If you are not using the skeleton code please give us detailed step-by-step instructions on how to setup and run your solution. (Tell us how you installed YOLO)

Milestone 3

Before you leave today, it is recommended you take some photos of the physical fruit from known distances so you can validate your pose estimation for the physical robot at home

Remember to submit your M3 before next lab session

Milestone 2 marking


We've looked over some of the results for M2, grades will be released when all groups have finished being marked off and looked over.

Some things we've noticed:

- Some physical maps have the correct shape but incorrect scale

You should check that your physical camera intrinsic parameters are correct

You should place an ArUco marker in front of the robot and see if it can correctly estimate the distance of the marker like how we validated the `update()` function in M2_3 slides



- Think about your use of time and driving strategy, some groups spent the entire 5 minutes on a single run while others completed 3 runs within the same time.

When the covariance in the markers get small, their estimate wont update that much